


The Upward-Spiral System

What was implemented, how the person-model changes over time, and how consent keeps reflection useful.



The central promise

Faerie does not try to finish describing a person. It keeps a versioned, correctable interpretation that can become narrower, less confident, or obsolete as life changes.

Implementation report - Phase 1 through Phase 6 | Verified focused suite: 266 passing tests | July 12, 2026

1. System architecture and authority


Each layer has one job. Facts, interpretations, chosen directions, and notification timing are deliberately separate so the model cannot turn a guess into identity or change a life plan without permission.

Layer	Allowed role	User authority	History
Memory	What happened or was explicitly said	Existing memory controls	Superseded facts remain
Person-model claim	Tentative pattern across evidence	Confirm, reject, or correct	Old wording and evidence remain
Investigation synthesis	Best current answer to one question	Approve, correct, reject, reopen	Every version is preserved
Suggested Investigation	Possibly useful uncertainty	Start, refine, defer, reject, never ask	Decisions block repetition
Growth tree	Chosen direction and experiments	Owns every mutation	Archived nodes stay visible
GoalAI proposal	Possible tree adjustment	Separate approval required	Stale proposals cannot apply
Reflection cadence	When Faerie may seek attention	Act, snooze, ignore, stop topic	Metadata only

Core safety boundary

Model output is a proposal. Approval of an Investigation synthesis does not silently approve a person-model change or a Growth-tree mutation; those are separate decisions.

2. How Investigations change over time



A candidate never starts itself


Faerie ranks possible questions by relevance, uncertainty, expected usefulness, burden, and sensitivity. At most two are visible and active capacity is capped at five. Sensitive topics require explicit permission.

A synthesis is a working interpretation

Each version stores confidence, evidence, counterevidence, unknowns, possible experiments, what changed, and reopen conditions. Meaningful answers or experiment outcomes make a new draft due. The user approves before it can influence downstream understanding.

New evidence	Expected response
Supporting answer	Confidence may rise, with source attached
Contradictory answer	Confidence may fall; interpretation may narrow
Successful exception	Create a new version; preserve the earlier generalization
Unhelpful experiment	Draft a lower-confidence synthesis
Insufficient evidence	Say so; keep unknowns explicit

3. How the Growth tree grows - and prunes



Growth is not permanent accumulation

Roots are chosen life directions, Branches are strategies or capabilities, and Leaves are small experiments. Outcomes flow back as evidence. Reusable learning can be harvested upward, but a proposed change still requires approval.

Gardening proposal	Purpose
Rewrite	Update language to match the person's current meaning
Split	Separate a broad or conflicted direction
Merge	Combine overlapping directions
Pause	Keep it without active pressure
Archive	Move it into visible history
Attach evidence	Ground the node without changing its meaning
Leave unchanged	Record that the goal still fits

New evidence can trigger a relevance review. A quiet active goal gets a gentle check after 30 days; descendant activity resets the clock. Paused and archived goals stay quiet.

4. Outcomes close the learning loop

A Leaf is an experiment, not a moral test. Completing, attempting, avoiding, and abandoning can all reveal something useful.


Outcome field	Why it matters
Result	Completed, attempted, avoided, or abandoned
What happened	Factual evidence, separate from interpretation
Expected obstacle	Compares the prediction to real life
Surprise	Finds exceptions and missing variables
Helpfulness (1-5)	Bad advice must reduce confidence
Changed understanding	User-authored correction to the working model
Next adjustment	Creates an approval-gated next Leaf proposal

Propagation after an outcome

Destination	What is added	What is not allowed
Goal evidence	Structured outcome reference	Automatic goal rewrite
Memory	Encrypted factual outcome	Identity-level conclusion
Inference	Evidence on a relevant claim	Silent claim confirmation
Investigation	Synthesis readiness and counterevidence	Overwriting old synthesis
Growth tree	Proposed next Leaf	Automatic child creation

This is the upward turn: lived reality is allowed to disprove Faerie.

5. Reflection rhythm and consent



One shared gate

Investigation check-ins, inference reviews, newly graduated hypotheses, and GoalAI updates all use the same local cadence. They cannot each send their own weekly prompt. Contradictions rank above routine check-ins but never bypass the cap.

Default	Behavior
Global limit	At most one unsolicited reflection every 7 days
Quiet hours	21:00 through 08:00 local time
Backlog	At most 3 queued reflection subjects
Snooze	3, 6, 12, then 24 days; capped at 28
Ignore	30, 60, then 90 days of topic-level quiet
Never	Durably disables that prompt topic
Explicit /remind	Bypasses cadence because the user scheduled it

The cadence database stores no prompt body and no private answer - only kind, opaque subject key, trigger, priority, timestamps, state, and optional usefulness/burden ratings.

6. Phases delivered

Phase	Implemented capability	Primary safeguard
1	Versioned Investigation synthesis	Approval before downstream influence
2	Person-model reconciliation: support, contradict, narrow, retire, situational, change over time	Strong identity claims need stronger evidence
3	Suggested Investigation engine	No auto-start; sensitive permission; durable rejection
4	Tree relevance and gardening	Every mutation is approval-gated
5	Structured experiment outcomes	Bad advice lowers confidence; history remains
6	Shared cadence and longitudinal evaluation	Quiet hours, weekly cap, backlog, suppression

Synthetic multi-month journeys

Journey	Contract verified
Changed dream	Can be reconsidered without erasing old meaning
Mistaken fear hypothesis	Contradiction gets priority; weekly consent remains
Successful exception	New synthesis version narrows the old generalization
Sensitive rejection	Never-ask remains durable across future attempts
Repeated burden	Snooze and ignore intervals expand
Conflicting evidence	History and ability to change are preserved

Focused verification: 266 passing tests across inference, Investigations, GoalAI, goals, outcomes, scheduler behavior, UI bridges, cadence policy, and longitudinal journeys.

7. What remains to learn from real use

The architecture is complete enough to begin longitudinal product tuning. The remaining questions are human, not merely technical.

Question	What to watch
Is weekly too quiet or too frequent?	Prompt usefulness and burden ratings over several months
Are goal checks timely?	Whether 30 quiet days feels supportive or guilt-inducing
Are suggestions insightful?	Start/refine/defer/reject/never-ask ratios
Does Faerie change its mind fairly?	Whether users recognize the cited evidence and explanation
Do outcomes feel lightweight?	Completion rate and whether the form feels like homework
Does the tree remain legible?	Growth versus archive rate and usefulness of history

Known engineering constraints

The Growth UI still shares the large memory.html file, so modularization would reduce the blast radius of future interface changes. The cadence uses explicit subject keys rather than broad semantic similarity; Suggested Investigations retain their stronger topic-key rejection boundary. Real burden data should guide tuning rather than optimizing for engagement.

Recommended checkpoint

Use the system normally for several weeks. Then review: which prompt arrived, why it arrived, whether it was useful, whether it felt heavy, and what Faerie changed afterward. That is the feedback loop that should tune the next version.

Faerie should feel attentive, corrigible, and increasingly specific - never certain about a person merely because it has accumulated data.